

AgentID — Business Requirements Document

Product: AgentID — the trust layer for AI agents **Tagline:** TLS for agentic services **Owner:** Lakshya Karira
Status: Draft **Audience:** Prospective customers, partners, investors, and collaborators **Last updated:** July 2026

1. Executive summary

Companies are deploying AI agents that act on their behalf — reading data, calling tools, moving money, and increasingly transacting with other companies' agents. But these agents run on a security and trust model that was never designed for them: static, long-lived API keys.

The consequence is a growing, largely unmanaged business risk. When an agent is compromised or misbehaves, there is no fast way to shut off its access everywhere at once, no way to control what data it sends across a boundary, and no trustworthy record of what it did. As agents move from experiments into production — and into regulated and cross-company use — this gap becomes a blocker to adoption.

AgentID is the trust layer that closes this gap. It gives every agent a verifiable identity, removes long-lived credentials from agents entirely, enforces what each agent is allowed to do, lets an organization revoke a compromised agent across every connected tool in seconds, and produces an independently verifiable audit trail. The positioning is deliberately simple: **the communication protocols for agents (such as MCP and A2A) define how agents talk; AgentID defines whether they should trust each other and what they're allowed to do — the same role TLS plays for the web.**

The near-term wedge is instant deprovisioning — a capability every team running agents in production will need the first time something goes wrong. The long-term opportunity is to become the default trust infrastructure for the emerging agent economy.

2. Business problem

2.1 The situation

Adoption of AI agents inside businesses is accelerating. Agents are being given access to sensitive systems — payments, code, customer data, internal APIs — to act autonomously.

2.2 The problem

The way agents authenticate and are governed has not kept pace. Today an agent typically holds a static API key. This creates four concrete business problems:

- **No fast containment.** If an agent is compromised, revoking its access is manual and slow — often hours or days, tool by tool. In an autonomous system, that delay is a serious exposure.
- **No data control.** Agents send full data payloads to the tools and third parties they call, with no fine-grained control over what sensitive information crosses each boundary.

- **No accountability.** There is no independently verifiable record of what an agent did, which is a growing problem as regulation and enterprise procurement demand auditability for automated actions.
- **No cross-company trust.** As agents from different companies begin to interact, there is no established way for one organization to trust another's agent without fragile, tightly-coupled integrations.

2.3 Why it matters now

Three forces make this urgent rather than theoretical: - **Production deployment.** Agents are leaving the lab; the cost of a mistake is now real. - **Agent-as-a-Service.** Companies are handing sensitive data to third-party agent providers with only contracts as protection. - **Regulation.** Emerging AI regulation and existing compliance regimes increasingly require audit trails and data governance for automated decision-making.

3. Business objectives

Objective	Description
Establish the category	Become recognized as the trust layer for AI agents — the "TLS for agentic services."
Solve an urgent present-day pain	Deliver instant, provable agent deprovisioning as the initial wedge into customers.
Build a defensible position	Achieve adoption strong enough that the credential format becomes a de facto standard.
Enable enterprise adoption of agents	Remove the trust and compliance blockers that currently slow agent deployment.
Create a path to the agent economy	Position to power cross-organization agent trust as that market emerges.

4. Success criteria (business terms)

Early stage - A working product in the hands of real users, with instant deprovisioning demonstrated end to end. - A small set of design-partner customers actively using the product. - Evidence of willingness to pay from at least the first buyer conversations.

Growth stage - A repeatable base of paying customers among companies running agents in production. - Adoption by developers and, ideally, integration by one or more agent frameworks or tool vendors. - Expansion within customers from the initial operational buyer into security and compliance budgets.

Long term - Recognition as the default trust layer referenced across the agent ecosystem. - First cross-organization agent-trust deployments in production.

(All targets to be quantified once early buyer conversations establish realistic benchmarks.)

5. Stakeholders

Stakeholder	Interest
Platform / infrastructure teams	Central control of agents; fast revocation; no credential sprawl. Primary early buyer.
Security & compliance leaders	Audit trails, policy enforcement, data governance. Expansion buyer.
AI engineers / developers	Simple, low-friction way to secure agents without changing how they build. Adoption driver.
Agent-as-a-Service providers	Ability to prove trustworthiness and win enterprise clients. High-value segment.
Tool / platform vendors	Potential ecosystem partners who adopt the standard.
Regulators (indirect)	Assurance that automated actions are governed and auditable.

Primary buyer: platform/infrastructure teams at companies running agents in production. **Budget owner:** CTO or CISO depending on company size.

6. Scope

6.1 In scope

- A trust layer that provides verifiable agent identity, credential management, authorization, data-policy enforcement, instant revocation, and verifiable audit.
- Multiple deployment options spanning a trust-versus-effort spectrum, so customers of different sizes and risk profiles can adopt.
- An open core to drive developer adoption, with commercial offerings for enterprise assurance needs.

6.2 Out of scope

- AgentID is **not** a general security suite (not a firewall, SIEM, malware scanner, or DLP platform).
- AgentID does **not** build agent frameworks or the communication protocols themselves — it complements them.
- AgentID does **not** claim to control what a third party does with data after it legitimately receives it; it controls what data is shared in the first place and provides verifiable records.

7. Business model (direction)

AgentID follows an **open-core** model, consistent with its goal of becoming a standard:

- **Free and open** — the core developer tooling and single-organization enforcement, which drives adoption and answers the natural concern about trusting a proprietary security layer.

- **Commercial** — the hosted control plane, management dashboard, cross-organization federation, compliance-grade audit, premium deployment modes, and support.

Pricing direction is a tiered platform fee rather than per-call metering, which fits the deployment model (including self-hosting, where usage cannot be metered) and avoids taxing the adoption the business depends on. Specific pricing to be set following buyer discovery.

The principle: **free where it drives adoption, paid where it delivers enterprise assurance.**

8. Business risks & assumptions

Risk / assumption	Description	Mitigation / note
Timing risk	The strongest differentiator (cross-organization trust) may be less urgent than present-day pains.	Sell the acute present pain (deprovisioning, audit); architect for the future.
Platform risk	Communication protocols (MCP, A2A) or a major platform could absorb a basic trust layer.	Position as the neutral, cross-vendor layer no single platform wants to own; go deeper than a protocol would.
Adoption chicken-and-egg	Full value grows as more tools verify AgentID; early on, adoption is voluntary.	Deliver standalone value on the customer side first; pursue a flagship design partner.
Trust concern	Customers may hesitate to route agent traffic through a third party.	Offer self-hosting where AgentID never sees customer data; open-source the enforcement path.
Founder capacity	Early stage led by a solo, part-time founder.	Prioritize a shippable wedge; validate before scaling commitment.
Market maturity	The agent-security category is young and fast-moving.	Validate assumptions continuously through buyer conversations; revisit the plan as the market develops.

Key assumptions the business case depends on: that agents will continue moving into production at scale; that enterprises will pay for containment, control, and auditability of agent behaviour; and that no incumbent protocol will fully close the trust gap first.

9. High-level cost & benefit

Cost drivers (early stage): primarily founder time to build and validate the wedge product, minimal infrastructure, and the cost of early customer discovery. Deliberately lean until product-market fit signals appear.

Expected benefit: a defensible position in an emerging infrastructure category with the potential to become a standard — a category whose comparables (identity, authentication, and security infrastructure for the web) have produced durable, high-value businesses. The near-term benefit is solving a concrete, urgent customer pain; the long-term benefit is owning the trust layer for the agent economy.

10. Recommendation

Proceed with a deliberately narrow first step: build and validate the instant-deprovisioning wedge with a small set of design-partner customers, confirm willingness to pay, and use that traction as the foundation for the broader trust-layer vision. Defer heavier investments (cross-organization federation, premium deployment modes) until pulled by real customer demand. This sequences risk sensibly — proving the business case on a sharp, present-day pain before committing to the larger, longer-horizon opportunity.

Honesty note: figures and market specifics are to be validated through buyer discovery and current market research; this document states the business case and the assumptions it rests on, so that any assumption proven false can be recognized as affecting the case.